

Homework 2: Deep Learning

Out April 26; Due May 3, beginning of exercise session*

Kristian Kersting, Alejandro Molina, Jinseok Nam

{kersting, molina, nam}@cs.tu-darmstadt.de

upload link: <https://www.dropbox.com/request/4B0KMkPqUZMHv4uo0IKU>

0. Assuming that you have implemented a feed-forward neural network architecture with the following constraints.




Figure 1: A feed-forward neural network architecture consisting of two hidden layers.

- The sigmoid activation function is used for computing hidden activations.
- No activation function in the output layer.
- In this homework, please use the following code snippet to initialize all parameters, inputs and their ground truth in the FIRST cell of your notebook.

```
# Python CODE  
import numpy as np  
# DO NOT FORGET TO SPECIFY THE SAME SEED  
np.random.seed(12345)
```

*We will discuss the solutions in the exercise session. It is my suggestions that you try to address at least 50% of the exercise questions. Simply try hard to solve them. This way, you will get familiar with the technical terms and with the underlying ideas of the lecture.

```

def initialize(input_dim, hidden_dim, output_dim, batch_size):
    W1 = np.random.randn(hidden_dim, input_dim) * 0.01
    b1 = np.zeros((hidden_dim,))
    W2 = np.random.randn(hidden_dim, hidden_dim) * 0.01
    b2 = np.zeros((hidden_dim,))
    W3 = np.random.randn(output_dim, hidden_dim) * 0.01
    b3 = np.zeros((output_dim,))

    # list of all network parameters
    parameters = [W1, b1, W2, b2, W3, b3]

    # minibatch of input instances
    x = np.random.rand(input_dim, batch_size)

    # ground truths
    y = np.random.randn(output_dim, batch_size)

    return parameters, x, y

# initialize parameters, inputs and targets
parameters, x, y = initialize(3, 4, 2, 5)

```

1. Given the forward pass of the above network architecture discussed during the exercise slot on April 26, please implement a squared loss function as follows

$$\mathcal{L}(\theta; \mathbf{x}, \mathbf{y}) = \frac{1}{M} \sum_{n=1}^M \frac{1}{2} \|f_{NN}(\mathbf{x}_n) - \mathbf{y}_n\|^2 \quad (1)$$

where f_{NN} denotes the outputs of the network given inputs \mathbf{x} , \mathbf{y} is the targets, and θ denotes a set of the parameters.

2. Implement a *backward* pass of the network to propagate errors calculated in eq. (1).

3. Write a *gradient checking* program to make sure that you have implemented the backward pass correctly.

HINT: You need to compute the numerical gradients of the loss function. For further details, see the lecture slides.

4. Please submit your notebook file that shows the followings at the end:

- the gradients of the loss function with respect to each parameter separately
 $\frac{\partial \mathcal{L}}{\partial \mathbf{W}_1}, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_1}, \frac{\partial \mathcal{L}}{\partial \mathbf{W}_2}, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_2}, \frac{\partial \mathcal{L}}{\partial \mathbf{W}_3}, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_3}$
- output of the gradient checking program. (it should be very small real value)